

Lab 4

Multi-tasking and Scheduler activations

Course: Operating Systems

April 13, 2025

Goal This lab helps the student practice scheduler activation in multi-tasking context (process or thread based implementation), and figure out why we need the multi-task management framework.

Contents In detail, this lab requires the student practice experiments using scheduler activation to provide the multi-tasking environment:

- CPU scheduler
- Dispatcher
- Scheduler policy description (using crontab)

Besides, the practices include the implementation of self-setup multi-tasking framework called **bktpool** (BK task pool). In addition to support implicit threading technique, we might cover further model of creation and management of threads i.e. fork-join, OpenMP (or CUDA), Grand Central Dispatch, Thread Building Block etc.

Result After doing this lab, student can understand the framework of multi-tasking using the techniques above to provide the scheduling feature.

Requirements Student need to review the theory of multi-tasking and scheduling.

Contents

1	Background	3
1.1	Multi-tasking environment	3
1.2	Scheduling subsystem	4
2	Programming Interfaces	6
2.1	Multi-task programming interface	6
2.1.1	fork() API	6
2.1.2	pthread_create API	6
2.1.3	clone() API	6
2.2	Signal handler API	6
2.2.1	kill and tkill	6
2.2.2	sigwait() API	7
2.2.3	sigaction() API	7
2.3	BK TaskPool API	7
2.3.1	Task declaration API	7
2.3.2	Task Pool Usage API	8
2.4	A Linux Scheduler - Cron	9
2.4.1	Introduction to Cron	9
2.4.2	Crontab	10
2.4.3	Crontab Syntax and Operators	11
2.4.4	Linux Crontab Command	12
3	Practices	13
3.1	Multitasking framework illustration BK_TPool	13
3.1.1	Create a set of resource entities	13
3.1.2	CPU scheduler	14
3.1.3	Dispatcher	14
3.1.4	Finalize task pool and resource worker	15
3.2	Practice with CronTab	18
4	Exercise	20
A	Appendix - Invoke System Call in UNIX-like system	21
A.1	Obtain system call list	21
A.2	Invoke system call using syscall() function	23
A.3	An example with time-related system call	23
B	Appendix - Invoke System Call in NT system	25
B.1	Obtain preprocessed header	25
B.2	Obtain preprocessed header	25
B.3	An example with time-related system call in NT	26

1 Background

1.1 Multi-tasking environment

Multicore or multiprocessor systems putting pressure on programmers, challenges include:

- Dividing activities
- Balance
- Data splitting
- Data dependency
- Testing and debugging

The task is an abstract entity to quantize the CPU computation power. We can implement it using the both concepts introduced in the first few chapter of Operating System course include process creating with fork system call or thread creating with thread library such as POSIX Thread (aka pthread). Despite of the comfortable of using the provided library, thread has a long history of development from userspace (down) to kernel space.

When the thread are placed in userspace or the legacy code, the mapping model is N:1 in which multiple thread is mapped in to one kernel thread and the scheduler has to decide which user thread are dispatch to take owner the computation resource. In this work, we deal with the same problem as the legacy thread library. We develop a scheduling subsystem to deploy multi-task on top a limit hardware computation resource.

There are some other concerned approaches in multi-tasking framework development:

Control using Signals are used in UNIX systems to notify a process that a particular event has occurred.

Communication using Shared memory or Message passing

Resouce sharing management Scheduling subsystem

Multithreading model

Many-to-one Many user-level threads mapped to single kernel thread

- One thread blocking causes all to block
- Multiple threads may not run in parallel on multicore system.
- Example system GNU Portable Thread (Few system currently use this model)

One-to-one Each user-level thread maps to one kernel thread

- Creating a user-level thread creates a kernel thread
- Number of threads per process sometimes restricted due to overhead.
- Example systems: Window, Linux

Mnay-to-many Allows many user level threads to be mapped to many kernel threads

- Allows the operating system to create a sufficient number of kernel threads
- Example system: Windows with the ThreadFiber package (Otherwise not very common)
- Example systems: Window, Linux

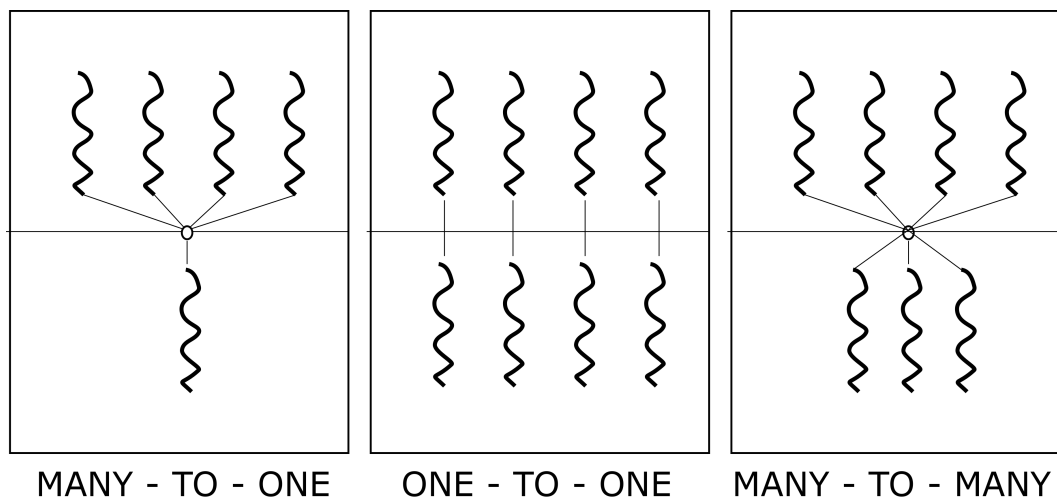


Figure 1: Multi-threading Model

The thread issues include:

- Semantics of `fork()` and `exec()` system calls
- Signal handling
 - Synchronous and asynchronous
- Thread cancellation of target thread
 - Asynchronous or deferred
- Thread-local storage
- Scheduler Activations

1.2 Scheduling subsystem

The CPU scheduler selects one process from among the processes in ready queue, and allocates the CPU core to it

Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:

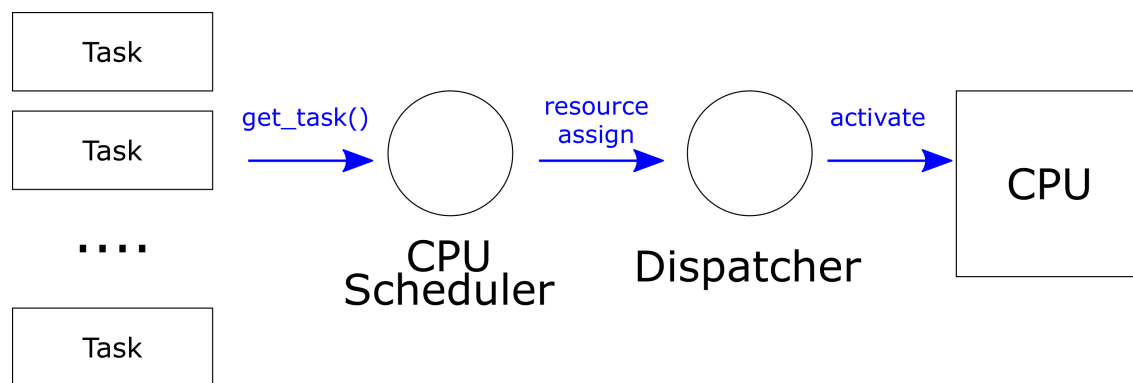


Figure 2: The two components of scheduling system

2 Programming Interfaces

2.1 Multi-task programming interface

2.1.1 fork() API

creates a new process by duplicating the calling process. The new process, referred to as the child, is an exact duplicate of the calling process, referred to as the parent

```
#include <unistd.h>
pid_t fork(void)
```

2.1.2 pthread_create API

creates a new thread start by a predeclared function in the calling process.

```
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,
                  void *(*start_routine) (void *), void *arg);
```

Notice: Compile and link with -pthread.

2.1.3 clone() API

The system call is the backend for both fork() API and pthread_create() API. clone() creates a new process, in a manner similar to fork(2). It is actually a library function layered on top of the underlying clone() system call. The superior of system call clone is the backend to provide a thread creation inside a thread. Pthread_create() is so-called a wrapper of system call clone().

(From the clone user manual)

CLONE_THREAD (since Linux 2.4.0-test8) If CLONE_THREAD is set, the child is placed in the same thread group as the calling process. To make the remainder of the discussion of CLONE_THREAD more readable, the term "thread" is used to refer to the processes within a thread group. Thread groups were a feature added in Linux 2.4 to support the POSIX threads notion of a set of threads that share a single PID. Internally, this shared PID is the so-called thread group identifier (TGID) for the thread group.

```
#define _GNU_SOURCE
#include <sched.h>

int clone(int (*fn)(void *), void *stack, int flags, void *arg, ...
         /* pid_t *parent_tid, void *tls, pid_t *child_tid */ )
```

2.2 Signal handler API

2.2.1 kill and tkill

The set of system call support sending a signal to a thread or process

The **kill()** system call can be used to send any signal to any process group or process.

The **tgkill()** system call sends the signal **sig** to the thread with the thread ID **tid** in the thread group **tgid**. The **kill()** is an obsolete predecessor to **tgkill()**. It allows only the target thread ID to be specified, which may result in the wrong thread being signaled if a thread terminates and its thread ID is recycled. Avoid using this system call.

```
#include <signal.h>

int kill(pid_t pid, int sig);
-----
#include <signal.h>          /* Definition of SIG* constants */
#include <sys/syscall.h>      /* Definition of SYS_* constants */
#include <unistd.h>

int syscall(SYS_tkill, pid_t tid, int sig);
-----
#include <signal.h>

int tgkill(pid_t tgid, pid_t tid, int sig);
```

2.2.2 sigwait() API

The function suspends execution of the calling thread until one of the signals specified in the signal set **set** becomes pending. The function accepts the signal (removes it from the pending list of signals), and returns the signal number in **sig**.

```
#include <signal.h>

int sigwait(const sigset_t *restrict set, int *restrict sig);
```

2.2.3 sigaction() API

The system call is used to change the action taken by a process on receipt of a specific signal.

```
#include <signal.h>

int sigaction(int signum, const struct sigaction *restrict act,
              struct sigaction *restrict oldact);
```

2.3 BK TaskPool API

2.3.1 Task declaration API

Task definition requires a job execution function. We share almost the same API with other library by defining **task_init** include the 2 information of what function task will be executed and the argument to passing to the function.

```
int bktask_init(int *taskid, void *(*start_routine) (void *), void *arg);
```

An example of new task declarations:

```
int func(void *arg)
{
    int id = *((int *) arg);

    printf("Task-func -- Hello from %d\n", id);
    fflush(stdout);

    return 0;
}

int main()
{
    ...
    id[0] = 1;  bktask_init(&tid[0], &func, (void*) &id[0]);
    id[1] = 2;  bktask_init(&tid[0], &func, (void*) &id[1]);
    id[2] = 5;  bktask_init(&tid[0], &func, (void*) &id[2]);
    ...
}
```

2.3.2 Task Pool Usage API

Defined task is passed to assigned worker and is dispatched to be executed.

```
#include "bktpool.h"

int bkwrk_get_worker();
-----
int bktask_assign_worker(int bktaskid, int wrkid);
-----
int bkwrk_dispatch_worker(int wrkid);
```

An example of using Task Pool API

```
int main()
{
    ...
    wid[1] = bkwrk_get_worker();
    ret = bktask_assign_worker(tid[0], wid[1]);
    if (ret != 0)
        printf("assign_task_failed - tid=%d - wid=%d\n", tid[0], wid[1]);

    bkwrk_dispatch_worker(wid[1]);
    ...
}
```





Figure 3: The calling procedure of BK Task Pool's routines

2.4 A Linux Scheduler - Cron

2.4.1 Introduction to Cron

Cron is a scheduling daemon that executes tasks at specified intervals. These tasks are called cron jobs and are mostly used to automate system maintenance or administration.

For example, we can set up a cron job to automate repetitive tasks such as backing up databases or data, updating the system with the latest security patches, checking disk space usage, sending emails, etc.

The cron jobs can be scheduled to run by a minute, hour, day of the month, month, day of the week, or any combination of these. For Ubuntu, we can inspect the cron service by running the following command:

```
$ sudo systemctl status cron.service
```

In the case there is an error of "command not found", you need to check and install the package from the repository

```
# Verify the existence of the required package
$ sudo dpkg -l | grep systemd
ii  libpam-systemd:amd64      system and service manager - PAM module
ii  libsystemd-daemon0:amd64 systemd utility library
ii  libsystemd-login0:amd64  systemd login utility library
ii  systemd-services         systemd runtime services
ii  systemd-shim             shim for systemd
```

```
# Install the package "systemd"
```

```
$ sudo apt-get install systemd
```

```
# When the package is installed, the command is working now:
```

```
$ sudo systemctl status cron.service
```

```
cron.service
```

```
Loaded: error (Reason: No such file or directory)
```

```
Active: inactive (dead)
```

Cron Cycle on Linux

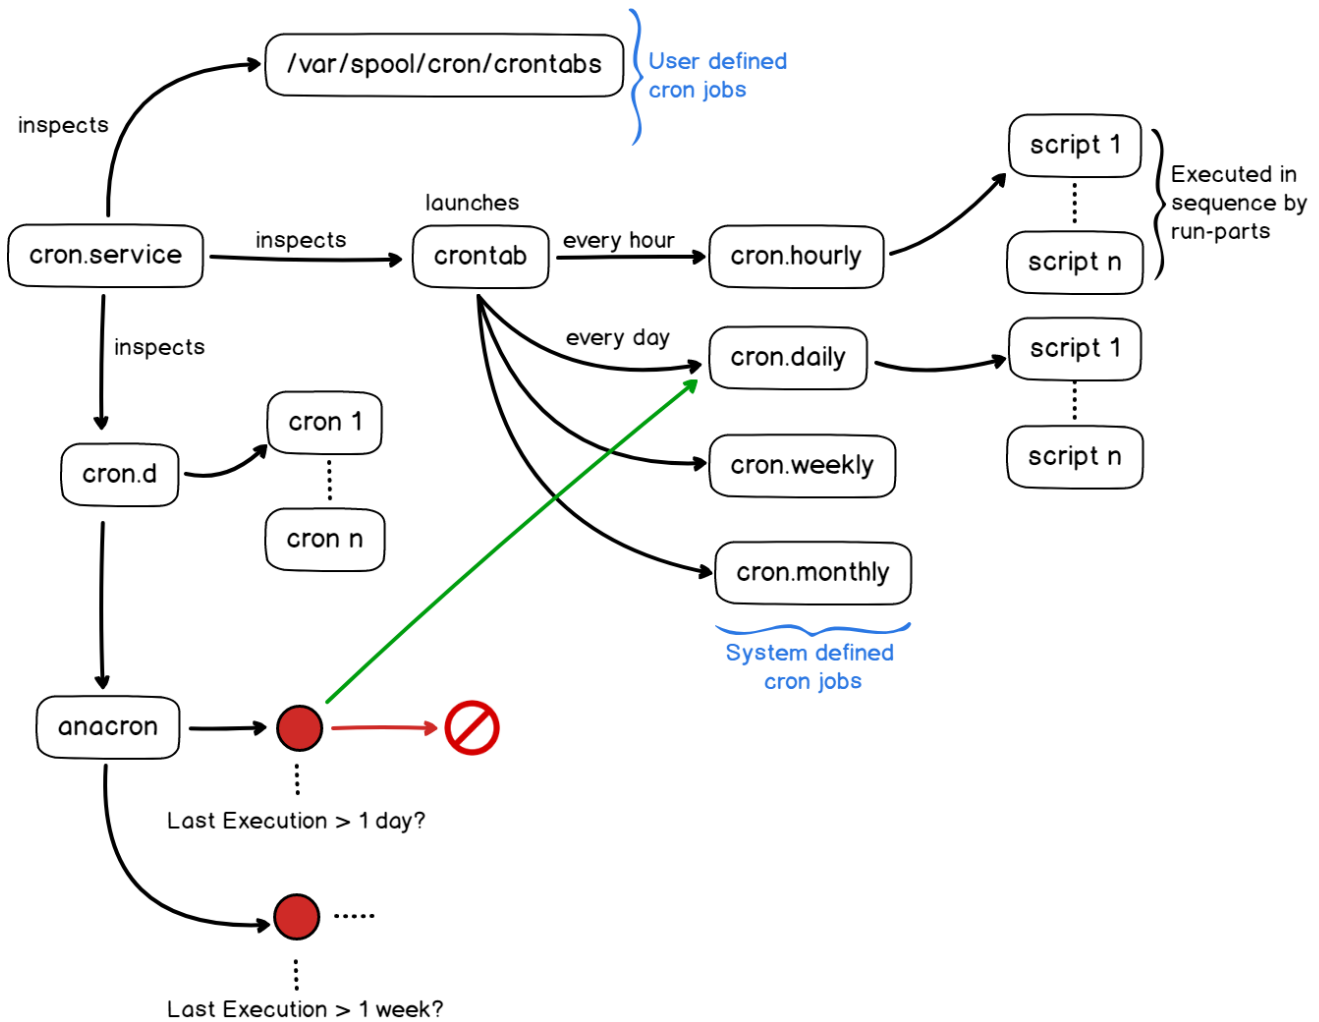


Figure 4: Cron Cycle

2.4.2 Crontab

Crontab (cron table) is a text file that specifies the schedule of cron jobs. There are two types of crontab files. The system-wide crontab files and individual user crontab files. Users' crontab files are named according to the user's name, and their location is at the `/var/spool/cron/crontabs` directory on Ubuntu. The `/etc/crontab` file and the scripts inside the `/etc/cron.d` directory are system-wide crontab files that can be edited only by the system administrators.

Figure 4 describes the cron cycle. Cron will inspect the user defined cron jobs and execute them if needed. It will also inspect the crontab file where several default cron jobs are defined by default.

Those default cron jobs are scripts that instructs your host to verify every minute, every hour, every day and every week specific folders and to execute the scripts that are inside them.

Finally, the `cron.d` directory is inspected. The `cron.d` may contain custom cron files and it also contains a very important file which is the `anacron` cron file.

2.4.3 Crontab Syntax and Operators

Each line in the user crontab file contains six fields separated by a space followed by the command to be run:

* * * * *	command(s)	output
- - - - -		
- - - -	Day of week (0 - 7) (Sunday=0 or 7) (Mon - Sun)	
- - - - -	Month (1 - 12)	
- - - - - - -	Day of month (1 - 31)	
- - - - - - - -	Hour (0 - 23)	
- - - - - - - - -	Minute (0 - 59)	

The detail of a cron command are:

- The first five fields * * * * * specify the time/date and recurrence of the job.
- In the second section, the command specifies the location and script you want to run.
- The final segment output is optional. It defines how the system notifies the user of the job completion. Cron will issue an email if there is any output from the cron job. Cron jobs are meant to not produce any output if everything is ok.

The first five fields may contain one or more values, separated by a comma or a range of values separated by a hyphen.

- “*” - The asterisk operator means any value or always. If you have the asterisk symbol in the Hour field, it means the task will be performed each hour.
- “,” - The comma operator allows you to specify a list of values for repetition. For example, if you have 1,3,5 in the Hour field, the task will run at 1 am, 3 am, and 5 am.
- “-” - The hyphen operator allows you to specify a range of values. If you have 1-5 in the Day of the week field, the task will run every weekday (From Monday to Friday).
- “/” - The slash operator allows you to specify values that will be repeated over a certain interval between them. For example, if you have */4 in the Hour field, it means the action will be performed every four hours. It is same as specifying 0,4,8,12,16,20. Instead of an asterisk before the slash operator, you can also use a range of values, 1-30/10 means the same as 1,11,21.

There are several special Cron schedule macros used to specify common intervals. We can use these shortcuts in place of the five-column date specification:

- *@yearly* (or *@annually*) - Run the specified task once a year at midnight (12:00 am) of the 1st of January. Equivalent to 0 0 1 1 *.

- *@monthly* - Run the specified task once a month at midnight on the first day of the month. Equivalent to `0 0 1 * *`.
- *@weekly* - Run the specified task once a week at midnight on Sunday. Equivalent to `0 0 * * 0`.
- *@daily* - Run the specified task once a day at midnight. Equivalent to `0 0 * * *`.
- *@hourly* - Run the specified task once an hour at the beginning of the hour. Equivalent to `0 * * * *`.
- *@reboot* - Run the specified task at the system startup (boot-time).

Example:

1. `10 10 1 * * /path/to/script.sh`: A cron job that executes every 10:10 AM each month on the first day.
2. `23 0-23/2 * * * /path/to/scripts.sh`: Execute the script.sh at 23 minutes after midnight, 2 am, 4 am..., everyday

Exercise: Convert the following intervals to crontab presentation:

- Every Monday at 08:30
- Every workday, every 30 minutes, from 8:15 to 17:45
- Last day of every month at 17:30

2.4.4 Linux Crontab Command

The crontab command allows you to install, view, or open a crontab file for editing:

- `crontab -e` - Edit crontab file, or create one if it doesn't already exist.
- `crontab -l` - Display crontab file contents.
- `crontab -r` - Remove your current crontab file.
- `crontab -i` - Remove your current crontab file with a prompt before removal.
- `crontab -u <username>` - Edit other user crontab file. This option requires system administrator privileges.

The cron daemon automatically sets several environment variables::

- The default path is set to `PATH=/usr/bin:/bin`. If the command you are executing is not present in the cron-specified path, you can either use the absolute path to the command or change the cron `PATH` variable. You can't implicitly append `PATH` as you would do with a regular script.
- The default shell is set to `/bin/sh`. To change to a different shell, use the `SHELL` variable.
- Cron invokes the command from the user's home directory. The `HOME` variable can be set in the crontab.

3 Practices

In this section, we work on step by step building up a multi-task framework.

3.1 Multitasking framework illustration BK_TPool

The task pool implement follows the below steps:

3.1.1 Create a set of resource entities

```
#include <signal.h>
#include <stdio.h>

#define _GNU_SOURCE
#include <linux/sched.h>
#include <sys/syscall.h>      /* Definition of SYS_* constants */
#include <unistd.h>
#define INFO
#define WORK_THREAD

int bkwrk_create_worker()
{
    unsigned int i;

    for (i = 0; i < MAX_WORKER; i++)
    {
#ifdef WORK_THREAD
        void **child_stack = (void **) malloc(STACK_SIZE);
        unsigned int wrkid = i;
        pthread_t threadid;

        sigset_t set;
        int s;

        sigemptyset(&set);
        sigaddset(&set, SIGQUIT);
        sigaddset(&set, SIGUSR1);
        sigprocmask(SIG_BLOCK, &set, NULL);

        /* Stack grow down - start at top */
        void *stack_top = child_stack + STACK_SIZE;

        wrkid_tid[i] = clone(&bkwrk_worker, stack_top,
                            CLONE_VM|CLONE_FILES,
                            (void *) &i);
#endif
#ifdef INFO
        fprintf(stderr, "bkwrk_create_worker - got worker-%u\n", wrkid_tid[i]);
#endif

        usleep(100);
    }
#else
    int bkwrk_create_worker()
    {
        unsigned int i;

        for (i = 0; i < MAX_WORKER; i++)
        {
#ifdef WORK_THREAD
            void **child_stack = (void **) malloc(STACK_SIZE);
            unsigned int wrkid = i;
            pthread_t threadid;

            sigset_t set;
            int s;

            sigemptyset(&set);
            sigaddset(&set, SIGQUIT);
            sigaddset(&set, SIGUSR1);
            sigprocmask(SIG_BLOCK, &set, NULL);

            /* Stack grow down - start at top */
            void *stack_top = child_stack + STACK_SIZE;

            wrkid_tid[i] = clone(&bkwrk_worker, stack_top,
                                CLONE_VM|CLONE_FILES,
                                (void *) &i);
#endif
#ifdef INFO
            fprintf(stderr, "bkwrk_create_worker - got worker-%u\n", wrkid_tid[i]);
#endif
        }
    }
#endif
}
```

```

    usleep(100);
#else

```

Step 3.1.1 Create resource instance using thread or process technique. The two kinds of instance can be initialized using the system call `clone()` or the wrapped library function `fork()` and `pthread_create()`.

Step 3.1.2 Set up the control signal masking with allowance of the two signal `SIGQUIT` or `SIGUSR1`.

3.1.2 CPU scheduler

```

int bkwrk_get_worker()
{
    wrkid_busy[1] != 0;

    return 1;

    /* TODO Implement the scheduler to select the resource entity */
}

```

3.1.3 Dispatcher

Assign worker a assign a task to a worker

```

int bktask_assign_worker(unsigned int bktaskid, unsigned int wrkid)
{
    if (wrkid < 0 || wrkid > MAXWORKER)
        return -1;

    struct bktask_t *tsk = bktask_get_byid(bktaskid);

    if (tsk == NULL)
        return -1;

    /* Advertise I AM WORKING */
    wrkid_busy[wrkid] = 1;

    worker[wrkid].func = tsk->func;
    worker[wrkid].arg = tsk->arg;
    worker[wrkid].bktaskid = bktaskid;

    printf("Assign tsk-%d wrk-%d\n", tsk->bktaskid, wrkid);
    return 0;
}

```

Assign worker a assign a task to a worker

```
int bkwrk_dispatch_worker(unsigned int wrkid)
{
#ifdef d_busy[wrkid_cur] != 0kORK_THREAD
    unsigned int tid = wrkid_tid[wrkid];

    /* Invalid task */
    if(worker[wrkid].func == NULL)
        return -1;

#ifdef DEBUG
    fprintf(stderr, "brkwrk-dispatch-wrkid-%d--send-signal-%u-\n", wrkid, tid);
#endif

    syscall(SYS_tkill, tid, SIG_DISPATCH);
#else
    /* TODO: Implement fork version to signal worker process here */
#endif
}
```

3.1.4 Finalize task pool and resource worker

Task pool data structure delaration and pool initialization function

```
/*
 * From bktpool.h
 */

#include <stdlib.h>
#include <pthread.h>

#define MAX_WORKER 10

#define WRK_THREAD 1
#define STACK_SIZE 4096

#define SIG_DISPATCH SIGUSR1

typedef void (*thread_func_t)(void *);

/* Task ID is unique non-decreasing integer */
int taskid_seed;

int wrkid_tid[MAX_WORKER];
```

```

int wrkid_busy[MAXWORKER];
int wrkid_cur ;

struct bktask_t{
    void (*func)(void * arg);
    void *arg;
    unsigned int bktaskid;
    struct bktask_t *tnext;
} *bktask;

int bktask_sz;

struct bkworker_t {
    void (*func)(void * arg);
    void *arg;
    unsigned int wrkid;
    unsigned int bktaskid;
};

struct bkworker_t worker[MAXWORKER];

/*
 * From bktpool.c
 */

#include "bktpool.h"

int bktpool_init()
{
    return bkwrk_create_worker();
}

```

Resource worker Take a loop of waiting for incoming control signal and do its job. After finishing the task work, it backs to waiting state to catch the next event.

```

void * bkwrk_worker(void * arg)
{
    sigset_t set;
    int sig;
    int s;
    int i = *((int *) arg); // Default arg is integer of workid
    struct bkworker_t *wrk = &worker[i];

    /* Taking the mask for waking up */

```



```

sigemptyset(&set);
sigaddset(&set, SIGUSR1);
sigaddset(&set, SIGQUIT);

#ifdef DEBUG
    fprintf(stderr, "worker-%i-start-living-tid-%d-\n", i, getpid());
    fflush(stderr);
#endif

while(1)
{
    /* wait for signal */
    s = sigwait(&set, &sig);
    if (s != 0)
        continue;

#ifdef INFO
    fprintf(stderr, "worker-wake-%d-up\n", i);
#endif

    /* Busy running */
    if (wrk->func != NULL)
        wrk->func(wrk->arg);

    /* Advertise I DONE WORKING */
    wrkid_busy[i] = 0;
    worker[i].func = NULL;
    worker[i].arg = NULL;
    worker[i].bktaskid = -1;
}
}

```

Worker data structure declaration and pool initialization function

```

void * bkwrk_worker(void * arg)
{
    sigset_t set;
    int sig;
    int s;
    int i = *((int *) arg); // Default arg is integer of workid
    struct bkworker_t *wrk = &worker[i];

    /* Taking the mask for waking up */
    sigemptyset(&set);
    sigaddset(&set, SIGUSR1);

```

```

sigaddset(&set , SIGQUIT);

#ifdef DEBUG
    fprintf(stderr , "worker-%i-start-living-tid-%d\n" , i , getpid());
    fflush(stderr);
#endif

    while(1)
    {
        /* wait for signal */
        s = sigwait(&set , &sig);
        if (s != 0)
            continue;

#ifdef INFO
        fprintf(stderr , "worker-wake-%d-up\n" , i);
#endif

        /* Busy running */
        if(wrk->func != NULL)
            wrk->func(wrk->arg);

        /* Advertise I DONE WORKING */
        wrkid_busy[i] = 0;
        worker[i].func = NULL;
        worker[i].arg = NULL;
        worker[i].bktaskid = -1;
    }
}

```

3.2 Practice with CronTab

In this lab, we will look at an example of how to schedule a simple script with a cron job. First, we will create a script called `date-script.sh` in our `HOME` folder to print the system date and time and appends it to a file. The content of our script is shown below:

```
#!/bin/sh

$ echo $(date) >> date-out.txt
```

Don't forget to make the script executable by using the `chmod` command. Then we define our job in the CronTab. To open the crontab configuration file for the current user, enter the following command:

```
$ crontab -e
```

We can add any number of scheduled tasks, one per line. In this case, we want to make our job run every minute, so we will add the following command (please change the `user` of the absolute path to your Linux username):

```
*/1 * * * * /home/user/date-script.sh
```

Wait for some minutes and verify our cron job by checking the content of our output file:

```
$ cat date-out.txt
```

```
Thu 20 Oct 2022 04:02:01 PM UTC  
Thu 20 Oct 2022 04:03:01 PM UTC  
Thu 20 Oct 2022 04:04:02 PM UTC  
Thu 20 Oct 2022 04:05:02 PM UTC  
Thu 20 Oct 2022 04:06:01 PM UTC
```

Practice

Convert the following intervals to crontab presentation

- Every Monday at 08:30
- Every workday, every 30 minutes, from 8:15 to 17:45
- Last day of every month at 17:30

4 Exercise

PROBLEM 1 Implement the FIFO scheduler policy to `bkwrk_get_worker()` in section 3.1.2.

Expected TaskPool Output

```
$ ./mypool
bkwrk_create_worker got worker 7593
bkwrk_create_worker got worker 7594
bkwrk_create_worker got worker 7595
bkwrk_create_worker got worker 7596
bkwrk_create_worker got worker 7597
bkwrk_create_worker got worker 7598
bkwrk_create_worker got worker 7599
bkwrk_create_worker got worker 7600
bkwrk_create_worker got worker 7601
bkwrk_create_worker got worker 7602
Assign tsk 0 wrk 0
worker wake 0 up
Task func - Hello from 1
Assign tsk 1 wrk 0 >>>>>>>>> Activate asynchronously
Assign tsk 2 wrk 1 >>>>>>>>> Activate asynchronously
worker wake 0 up
Task func - Hello from 2
worker wake 1 up
Task func - Hello from 5
```

PROBLEM 2 In section 3.1.1 You are provided a thread based implementation of task worker in the function `bkwrk_create_worker()`. Try to implement another version of the worker using more common `fork()` API.

PROBLEM 3 Base on the provided material of multi-task programming and signal control, develop your own framework of Fork-Join in theory.

A Appendix - Invoke System Call in UNIX-like system

A.1 Obtain system call list

In this section, we create an entry program, named "**sysheader.c**", uses a macro define to include the syscall header file as in listing A.1. Based on the live linking we will obtain the system call list.

Listing 1: sysheader.c source code

```
#include <sys/syscall.h>

int main()
{
    return 0;
}
```

Generate the header include file searching result by leveraging **gcc** with preprocessing option "-H"

```
$ gcc -H sysheader.c
. /usr/include/sys/syscall.h
.. /usr/include/asm/unistd.h
... /usr/include/asm/unistd_32.h
... /usr/include/asm/unistd_64.h
.. /usr/include/bits/syscall.h
Multiple include guards may be useful for:
/usr/include/bits/syscall.h
```

Then we obtain the list of system call in unistd_32 or unistd_64.h

```
$ cat /usr/include/asm/unistd_32.h
#ifndef _ASM_X86_UNISTD_32_H
#define _ASM_X86_UNISTD_32_H 1

#define __NR_restart_syscall 0
#define __NR_exit 1
#define __NR_fork 2
#define __NR_read 3
#define __NR_write 4
#define __NR_open 5
#define __NR_close 6
#define __NR_waitpid 7
#define __NR_creat 8
#define __NR_link 9
#define __NR_unlink 10
#define __NR_execve 11
...
```

The list of system call for x86_64 is

```
$ cat /usr/include/asm/unistd_64.h
#ifndef _ASM_X86_UNISTD_64_H
#define _ASM_X86_UNISTD_64_H 1

#define __NR_read 0
```

```
#define __NR_write 1
#define __NR_open 2
#define __NR_close 3
#define __NR_stat 4
#define __NR_fstat 5
#define __NR_lstat 6
#define __NR_poll 7
#define __NR_lseek 8
#define __NR_mmap 9
#define __NR_mprotect 10
#define __NR_munmap 11
...
```

An alternative way to obtain the libc symbolic constant definition is (*please map the bit syscall.h path based on your include searching results*)

```
$ cat /usr/include/bits/syscall.h
/* Generated at libc build time from syscall list. */
/* The system call list corresponds to kernel 4.17. */

#ifndef _SYSCALL_H
# error "Never use <bits/syscall.h> directly; include <sys/syscall.h> instead."
#endif

....

#ifdef __NR__sysctl
# define SYS__sysctl __NR__sysctl
#endif

#ifdef __NR_accept
# define SYS_accept __NR_accept
#endif

....

#ifdef __NR_clock_gettime
# define SYS_clock_gettime __NR_clock_gettime
#endif

#ifdef __NR_clock_nanosleep
# define SYS_clock_nanosleep __NR_clock_nanosleep
#endif
....
```

We use these specified numbers to invoke system call in the next section.

A.2 Invoke system call using `syscall()` function

`syscall()`¹ is a small library function that invokes the system call whose assembly language interface has the specified number with the specified arguments.

Symbolic constants for system call numbers can be found in the header file `<sys/syscall.h>`

The return value is defined by the system call being invoked. In general, a 0 return value indicates success. A -1 return value indicates an error.

```
#include <sys/syscall.h>          /* Definition of SYS_* constants */
#include <unistd.h>

long syscall(long number, ...);
```

A.3 An example with time-related system call

After obtaining system call number, to obtain the prototype of system call implementation and interface specification, i.e. return value, argument types, we need to get it from kernel source code. This is an example of `clock_gettime` system call.

```
int clock_gettime(clockid_t, struct timespec *);
```

Apply the include file searching we obtain in `/usr/include/bits/time.h` the definition of various `CLOCK_*` modes.

```
$ cat /usr/include/bits/time.h
# define CLOCK_REALTIME          0
# define CLOCK_MONOTONIC        1
# define CLOCK_PROCESS_CPUTIME_ID 2
# define CLOCK_THREAD_CPUTIME_ID 3
# define CLOCK_MONOTONIC_RAW    4
# define CLOCK_REALTIME_COARSE  5
# define CLOCK_MONOTONIC_COARSE 6
# define CLOCK_BOOTTIME         7
...
```

The structure `timespec` definition is obtained by applying the same previous include path searching `/usr/include/bits/types/struct_timespec.h`

```
$ cat /usr/include/bits/types/struct_timespec.h
#ifndef _STRUCT_TIMESPEC
#define _STRUCT_TIMESPEC
...
struct timespec
{
    __time_t tv_sec;          /* Seconds. */
    __syscall_slong_t tv_nsec; /* Nanoseconds. */
};
#endif
```

¹<https://man7.org/linux/man-pages/man2/syscall.2.html>

Using the system call list obtained in section A.1, we can invoke the system call `SYS_clock_gettime` as in the following program named "`sysptime.c`".

Listing 2: `sysptime.c` source code

```
#include <stdio.h>
#include <time.h>
#include <string.h>
#include <sys/syscall.h>      /* Definition of SYS_* constants */
#include <unistd.h>

int main()
{
    struct timespec ts;
    char buffer[22];
    const int bufsize = 21;

    syscall(SYS_clock_gettime, CLOCK_REALTIME, &ts);

    strftime(buffer, bufsize, "%Y-%m-%d-%H:%M:%S", gmtime(&ts.tv_sec));

    printf("Date: %s\n", buffer);
}
```

Compile and execute the program:

```
$ gcc -o sysptime sysptime.c
$ ./sysptime
Date: 2025-03-27 16:46:27
$ date
Thu Mar 27 16:46:27 UTC 2025
```


B Appendix - Invoke System Call in NT system

B.1 Obtain preprocessed header

Since Windows is closed source software, we obtain the interface description through window API document at <https://learn.microsoft.com/en-us/windows/win32/api/>

The system call description of NtQuerySystemTime

Requirements NtQuerySystemTime function (winternl.h)

Requirement	Value
Target Platform	Windows
Header	winternl.h
Library	ntdll.lib
DLL	ntdll.dll

Syntax in C++

```
__kernel_entry NTSTATUS NtQuerySystemTime(  
    [out] PLARGE_INTEGER SystemTime  
);
```

B.2 Obtain preprocessed header

```
#include <iostream>  
using namespace std;  
  
#include <Windows.h>  
#include <winternl.h>  
#include <timezoneapi.h>  
  
int main()  
{  
    return 0;  
}
```

Obtain the header definition.

```
c:\ws>cc /EHsc /P nthead.cpp  
nthead.cpp  
  
# The preprocessed file create name "nthead.i"  
c:\ws> more nthead.i  
#line 1 "nthead.cpp"  
#line 1 "C:\\Program Files (x86)\\Windows Kits\\10\\include\\10.0.26100.0\\um\\  
    Windows.h"  
#line 1 "C:\\Program Files (x86)\\Windows Kits\\10\\include\\10.0.26100.0\\shared\\  
    winapifamil
```

```

#pragma warning(push)
#pragma warning(disable:4001)
#line 23 "C:\\Program Files (x86)\\Windows Kits\\10\\include\\10.0.26100.0\\shared
\\winapifamily.h"
#pragma once
#line 25 "C:\\Program Files (x86)\\Windows Kits\\10\\include\\10.0.26100.0\\shared
\\winapifamily.h"
...

```

B.3 An example with time-related system call in NT

Using the previous preprocessed header we can obtain the structure definition of FILETIME, SYSTEMTIME and convert time function FileTimeToSystemTime()

Listing 3: nttime.cpp source code

```

#include <iostream>
using namespace std;

#include <Windows.h>
#include <winternl.h>
#include <timezoneapi.h>

#pragma comment(lib, "ntdll")

// Function prototype for NTQuerySystemTime
typedef NTSTATUS (NTAPI* pNTQuerySystemTime)(
    PLARGE_INTEGER SysTime
);

int main()
{
    PLARGE_INTEGER SysTime;
    SysTime = (PLARGE_INTEGER) new(LARGE_INTEGER);
    FILETIME f;
    SYSTEMTIME s;

    HMODULE mod=LoadLibrary("ntdll.dll");

    if (mod == NULL){
        std::cout << "Module loading failed!\n";
        return -1;
    }

    FARPROC pNTQuerySystemTime = GetProcAddress(mod, "NtQuerySystemTime");

    NTSTATUS status = NtQuerySystemTime(SysTime);

    FreeLibrary(mod);

    if (NT_SUCCESS(status)) {
        f.dwLowDateTime = SysTime->LowPart;

```

```
f.dwHighDateTime = SysTime->HighPart;
FileTimeToSystemTime(&f, &s);
std::cout << "Query-time-successfully-SysDate="
           << s.wYear << s.wMonth << s.wDay << std::endl;
} else {
    std::cout << "Query-time-failed" << std::endl;
    return -1;
}
}
```

Compile and execute the program:

```
$ cc /EHsc nttime.cpp
/out:nttime.exe
nttime.obj

c:\ws>nttime.exe
Query time successfully SysDate=2025413
```

Revision History

Revision	Date	Author(s)	Description
1.0	03.15	PD Nguyen	Document created
...	
2	10.2022	LHT Hoang	Update lab content, practices and exercises
3	10.2023	PD Nguyen	Add BK_TPool and related description, update exercises
3.1	03.2025	PD Nguyen	Add Appendix Invoke System Call